

# Distributed Data Processing

UA.DETI.CBD

José Luis Oliveira / Carlos Costa



1

## Introduction

- ❖ Large-Scale Data Processing
  - aim to use 1000s of CPUs, but with simplified managing
- ❖ MapReduce
- ❖ Apache Hadoop
- ❖ To start:
  - <https://www.youtube.com/watch?v=9s-vSeWej1U>



2

- Big Data
- Distributed
- Transmission im high latency

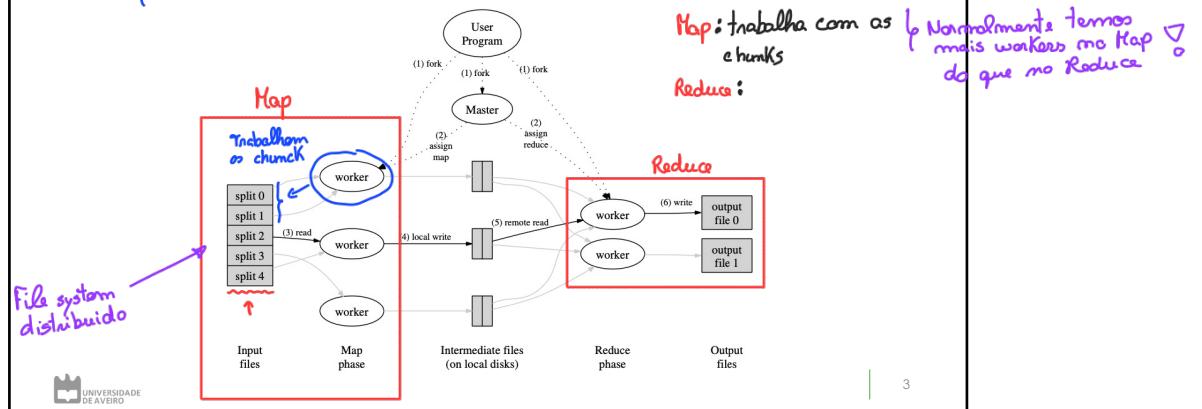
2

## What is MapReduce?

- "MapReduce is a programming model and an associated implementation for processing and generating large data sets"

→ • Dean, Jeffrey, and Sanjay Ghemawat. "MapReduce: Simplified data processing on large clusters." (2004). [Google]

AND GENERATING?



3

3

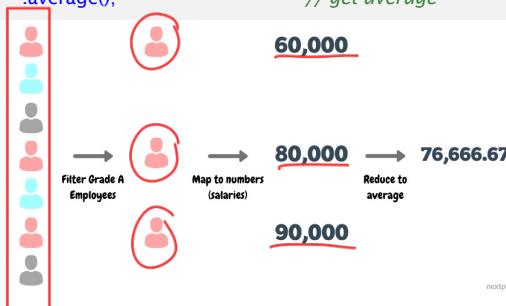
## What is MapReduce?

- Terms are borrowed from Functional Language (e.g., Lisp)
- Example: Sum of squares
  - **map** square '(1 2 3 4)  
Output: (1 4 9 16)  $y=x^2$ 
    - processes each record sequentially and independently
  - **reduce** + '(1 4 9 16) **Reduce**  
(+ 16 (+ 9 (+ 4 1)))  
Output: 30 **Sum**
    - processes set of all records in batches
- This concept has been reused in several programming languages and computational tools

## Examples: Java stream

```
List<String> values = Arrays.asList("1","2","3","4","5","6","7");
int soma = values.stream()
    .map(num->Integer.parseInt(num)) MapReduce
    .reduce(0, Integer::sum);
```

```
OptionalDouble avgSalary = employees.stream()
    .filter((e) -> e.grade == 'A')      // filter 'A' grade employees
    .mapToInt(e) -> e.salary)           // get IntStream of salaries
    .average();                         // get average
```



5

5

## Examples: MongoDB

```
Database:
db.orders.insertMany([
    { _id: 1, cust_id: "Ant O. Knee", ord_date: new Date("2020-03-01"), price: 25,
        items: [ { sku: "oranges", qty: 5, price: 2.5 }, { sku: "apples", qty: 5, price: 2.5 } ],
        status: "A" },
    { _id: 2, cust_id: "Ant O. Knee", ord_date: new Date("2020-03-08"), price: 70, items: [
        { sku: "oranges", qty: 8, price: 2.5 }, { sku: "chocolates", qty: 5, price: 10 } ],
        status: "A" },
    ...
])
Return the Total Price Per Customer:
var mapFunction1 = function() {
    emit(this.cust_id, this.price);
}; ↳ emit o par!
var reduceFunction1 = function(keyCustId, valuesPrices) {
    return Array.sum(valuesPrices);
};
db.orders.mapReduce(
    mapFunction1,
    reduceFunction1,
    { out: "map_reduce_example" }
)
db.map_reduce_example.find().sort( { _id: 1 } )
```

*Reduce*

{ "\_id": "Ant O. Knee", "value": 95 }  
{ "\_id": "Busby Bee", "value": 125 }  
{ "\_id": "Cam Elot", "value": 60 }  
...

<https://docs.mongodb.com/manual/tutorial/map-reduce-examples/>



6

6

## MapReduce frameworks

- ❖ Library/Tools that allow easily writing applications that process large amounts of data in parallel
  - distributed across several nodes
- ❖ Good retry/failure semantics
- ❖ Solution Pattern
  - many problems can be modelled in this way
- ❖ Implementations
  - **Hadoop**: the mapper and reducer are each a Java class that implements a particular interface.
  - **Spark**: newer and faster, it processes data in RAM using a concept known as an RDD, Resilient Distributed Dataset.



**Hadoop**: the mapper and reducer are each a Java class that implements a particular interface.

**Spark**: newer and faster, it processes data in RAM using a concept known as an RDD, Resilient Distributed Dataset.

↳ ele faz duplicação em vários nós !!

⇒ Tolerante a falhos!



7

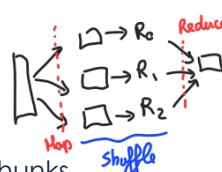
7

## MapReduce frameworks

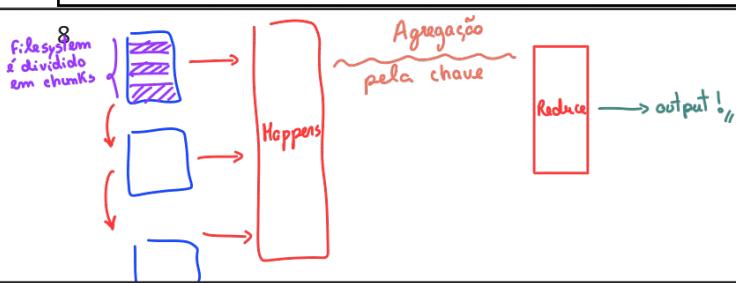
- ❖ Framework...
  - Automatic parallelization & distribution
  - Fault tolerance
  - I/O scheduling
  - Monitoring & status updates
- ❖ Two main tasks:
  - Mappers & Reducers
- ❖ Pipeline
  - Splits the input data-set into independent chunks
  - Mappers process of chunks in a parallel manner
  - aggregates the outputs of the maps
  - Reducers process the mappers output

• Ele faz:
 

- replicação
- balançamento de carga
- tolerância a falhos
- (...)



8



4

## MapReduce dataflow

### The map function

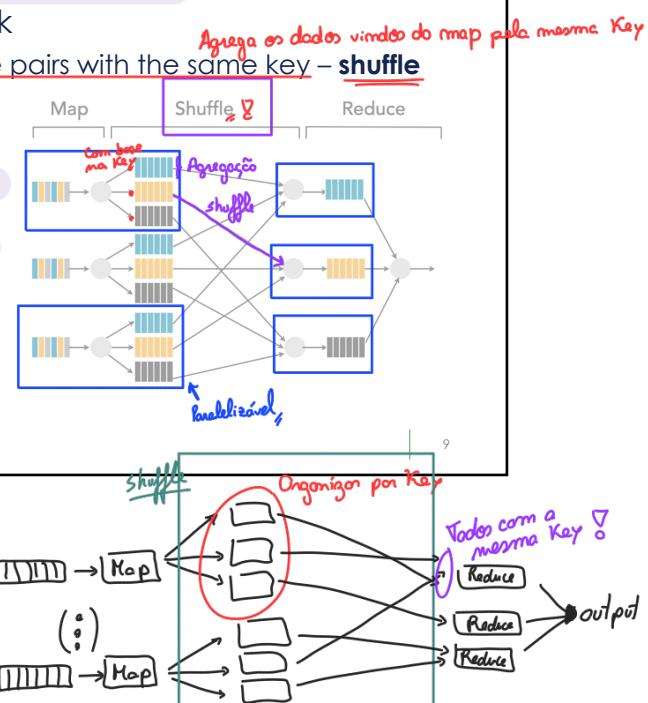
- is called once for every input record to extract (one or more) key-value from the input record

### MapReduce framework

- collects all the key-value pairs with the same key - **shuffle**

### The reduce function

- iterates over each collection of values with the same key and can produce output records
  - e.g., the number of occurrences of the key

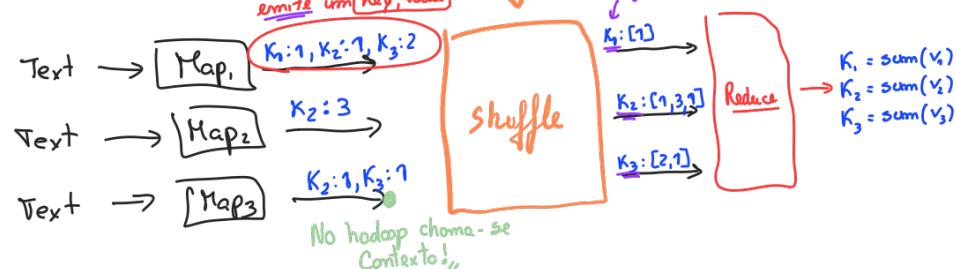
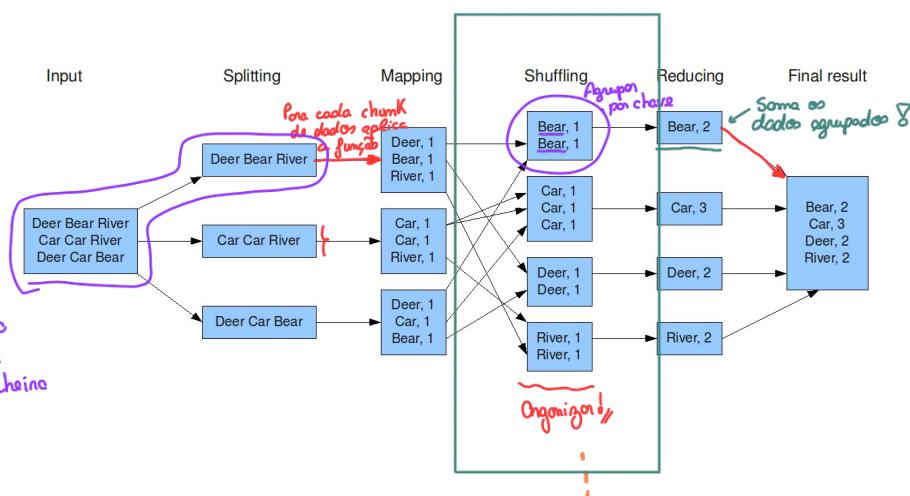


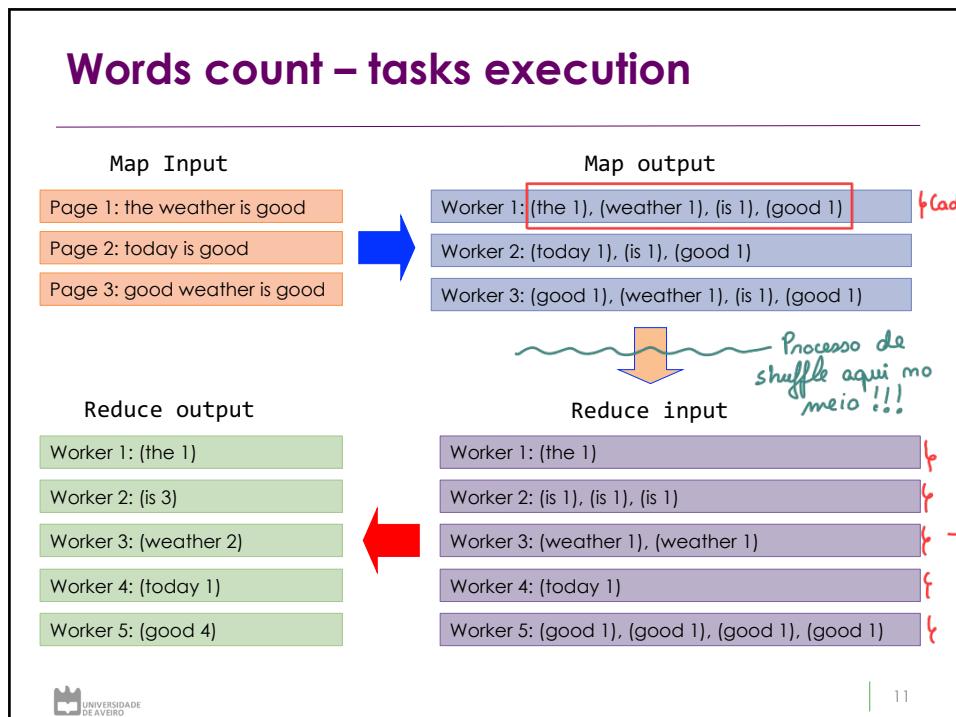
9

## Example: Words count dataflow

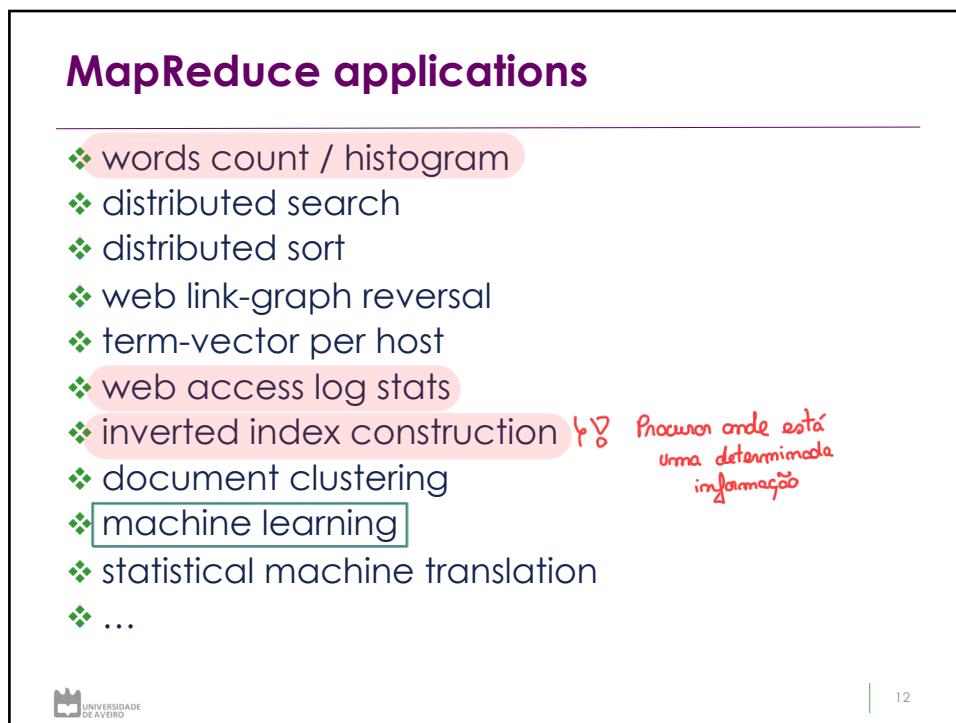
No hadoop ele fog por ficheiro quando é um diretório. Ou por limão quando é um ficheiro

10





11



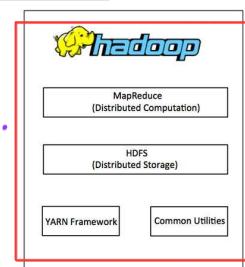
12

# Apache Hadoop

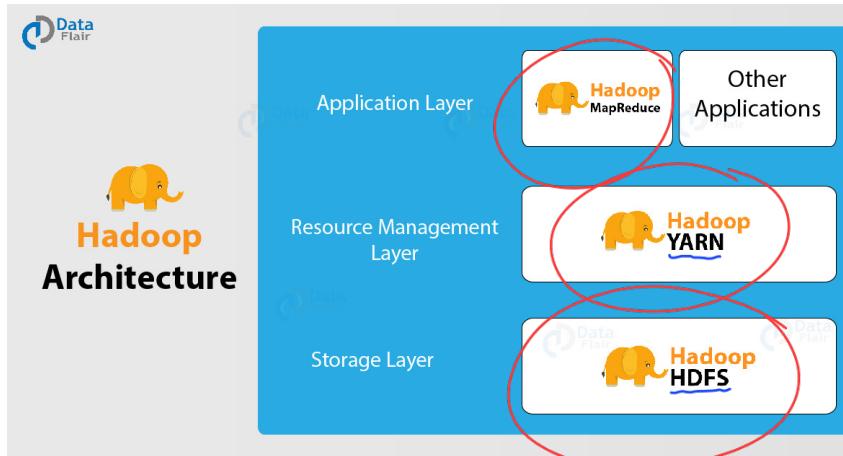


## Apache Hadoop Framework

- ❖ A framework that allows distributed processing of **large data sets across clusters of computers** using simple programming models.
- ❖ Open-source framework
  - Java
  - <https://hadoop.apache.org>
- ❖ **Main components**
  - Hadoop Distributed File System (**HDFS**)
    - Distributed, scalable, and portable file system
  - Hadoop Yet Another Resource Negotiator (**YARN**)
  - Hadoop Common Utilities
  - Hadoop **MapReduce**
    - Implementation of the MapReduce programming model



## Apache Hadoop Framework



15

15

## Hadoop Distributed File System (HDFS)

- ❖ Distributed file system designed to run on commodity hardware.
- ❖ Stores data redundantly on multiple nodes – Fault-tolerant file system
  - 3+ replicas for each block
  - Default Block Size : 128MB
- ❖ Master-Slave architecture
  - NameNode: Master tem os metadados e sabe localizar pelos slaves
  - DataNode: Slave onde estão guardados
- ❖ Typical usage pattern
  - Huge files (GB to TB)
  - Data is rarely updated
  - Reads and appends are common
    - Usually, random read/write operations are not performed

Hadoop  
Distributed  
File  
System

O YARN gera  
isto ...

NameNode : Master  
Data Node : Slave

{ é como um filesystem  
no nosso pc ...  
→ Nos são ficheiros virtuais  
fragmentados por n máquinas !!

Nome → chunk  
está guardado  
no Master !!

16

16

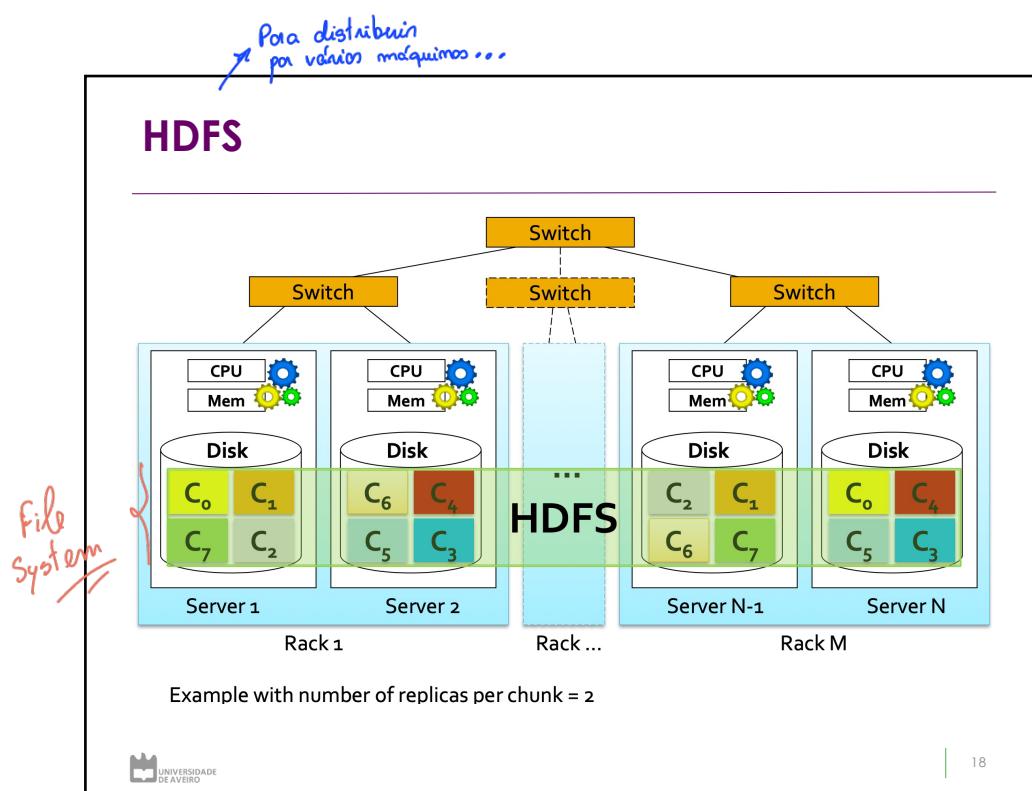
## HDFS – Assumptions and Goals

- ❖ **Hardware Failure**
  - An HDFS instance may consist of +100/+1000 of server machines.
- ❖ **Streaming Data Access**
  - HDFS is designed more for batch processing rather than interactive use. The emphasis is on high throughput of data access rather than low latency of data access.
- ❖ **Large Data Sets**
  - Applications that run on HDFS have large data sets. A typical file in HDFS is gigabytes to terabytes in size.
- ❖ **Simple Coherency Model**
  - HDFS applications need a write-once-read-many access model for files. A file once created, written, and closed need not be changed except for appends and truncates.
- ❖ **“Moving Computation is Cheaper than Moving Data”**
  - Especially true when the size of the data set is huge. Minimizes network congestion and increases the throughput of the system.
- ❖ **Portability Across Heterogeneous Platforms**

*Hámo com máquinas de diferentes sistemas Operativos*

17

17

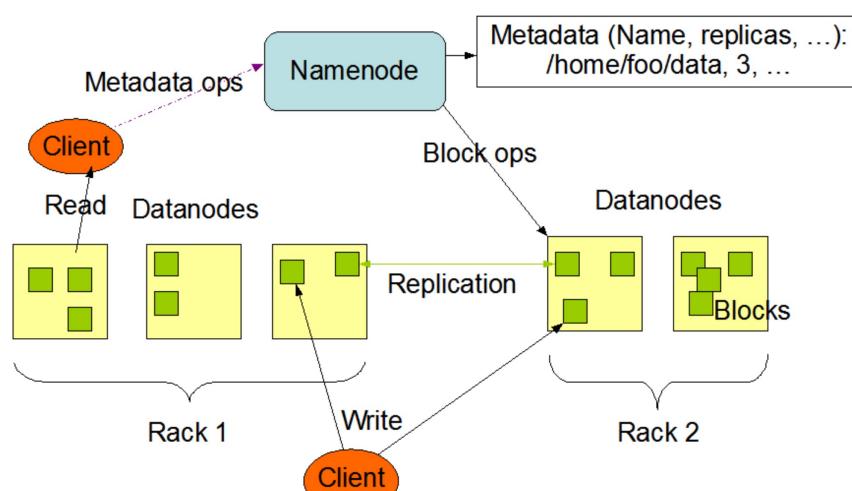


18

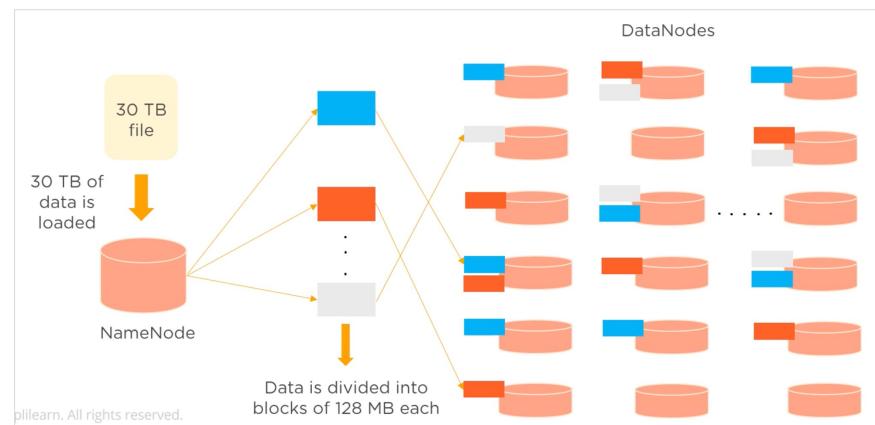
## HDFS

- ❖ **Master-Slave architecture**
  - Master: NameNode
  - Slave: DataNode
- ❖ **The Master node is a special node/server that**
  - Stores HDFS metadata
    - E.g., the mapping between the name of a file and the location of its chunks
  - Might be replicated
- ❖ Client applications: file access through HDFS APIs
  - Talk to the master node to find data/chunk servers associated with the file of interest
  - Connect to the selected chunk servers to access data

## Interaction between HDFS components



## Storage example



<https://www.youtube.com/watch?v=iANBytZ26MI>

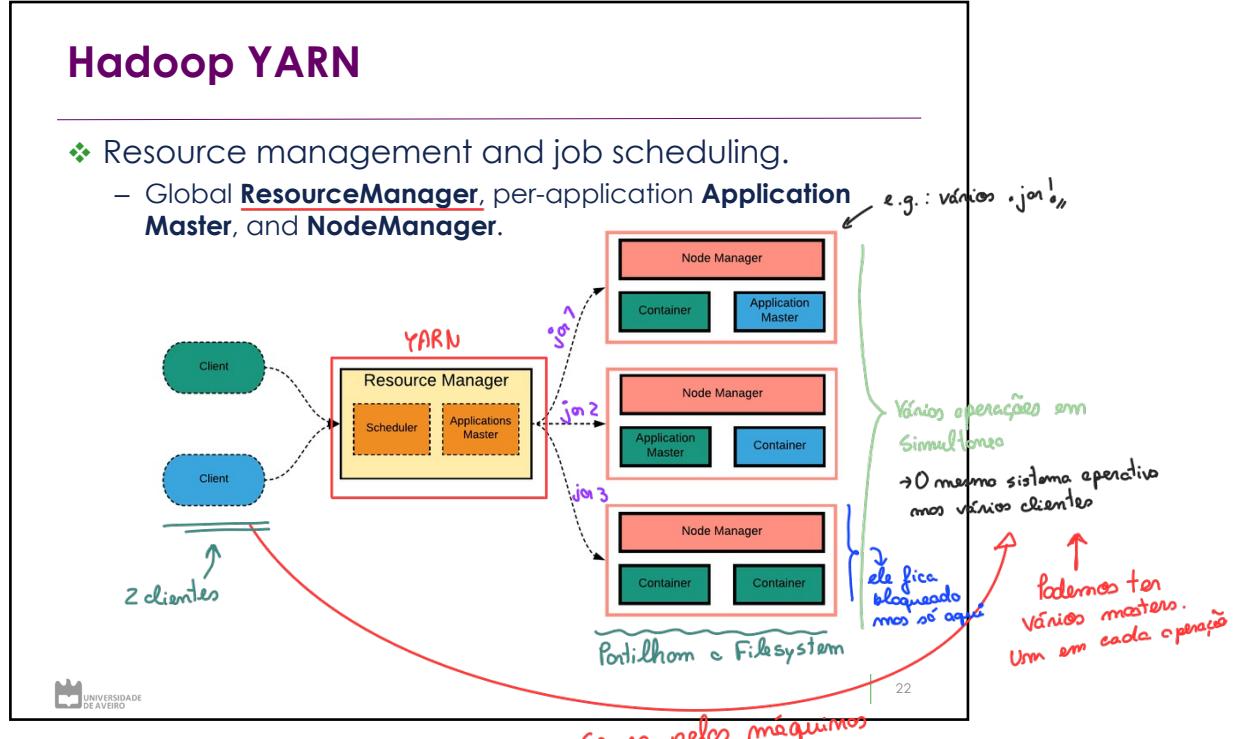
21

21

## Hadoop YARN

- Resource management and job scheduling.

- Global **ResourceManager**, per-application **Application Master**, and **NodeManager**.



22

## Hadoop MapReduce

- ❖ MapReduce is the data processing layer of Hadoop.
- ❖ MapReduce job comprises a number of map tasks and reduces tasks.
  - Each task works on a part of data. This distributes the load across the cluster.
  - **Map task** load, parse, transform and filter data.
  - Each **reduce task** works on the sub-set of output from the map tasks (e.g., by applying grouping and aggregation).



23

23

## Apache Hadoop

- ❖ Install & config
- ❖ Run
  - ~ [sbin ./start-dfs.sh](#)
- ❖ Basic help for all the commands
  - ~ [hadoop](#)
- ❖ Distributed file system commands
  - ~ [hadoop fs](#) ← *Poco file system*
- ❖ Execution of MapReduce jobs
  - ~ [hadoop jar](#) ← *Para executar o job!*



24

24

## HDFS commands

```

~ hadoop fs [generic options]
  [-appendToFile <localsrc> ... <dst>]
  [-cat [-ignoreCrc] <src> ...]
  [-checksum [-v] <src> ...]
  [-chgrp [-R] GROUP PATH...]
  [-chmod [-R] <MODE[,MODE]... | OCTALMODE> PATH...]
  [-chown [-R] [OWNER][:[GROUP]] PATH...]
  [-copyFromLocal [-f] [-p] [-l] [-d] [-t <thread count>] <localsrc> ... <dst>]
  [-copyToLocal [-f] [-p] [-ignoreCrc] [-crc] <src> ... <localdst>]
  [-count [-q] [-h] [-v] [-t <storage type>]] [-u] [-x] [-e] <path> ...]
  [-cp [-f] [-p | -p[topax]] [-d] <src> ... <dst>]
  [-createSnapshot <snapshotDir> [<snapshotName>]]
  [-deleteSnapshot <snapshotDir> <snapshotName>]
  [-df [-h] [<path> ...]]
  [-du [-s] [-h] [-v] [-x] <path> ...]
  [-expunge [-immediate] [-fs <path>]]
  [-find <path> ... <expression> ...]
  [-get [-f] [-p] [-ignoreCrc] [-crc] <src> ... <localdst>]
  ...
  ...

```



25

25

## Hadoop – HDFS Commands

*↳ quando estamos a usar clusters! //*

- ❖ Help
    - ~ `hadoop fs -help`
  - ❖ Creating a directory
    - ~ `hadoop fs -mkdir /myhdfsdir`
  - ❖ Listing a directory
    - ~ `hadoop fs -ls`
    - ~ `hadoop fs -ls /myhdfsdir`
  - ❖ Copy a file/dir from local file system to HDFS
    - ~ `hadoop fs -put /myhdfsdir/ /user/xpto/`
  - ❖ Get a file/dir from HDFS to local file system
    - ~ `hadoop fs -get /user/xpto/ /myhdfsdir/`
    - ... and many more *no disco local!*
- [https://www.tutorialspoint.com/hadoop/hadoop\\_command\\_reference.htm](https://www.tutorialspoint.com/hadoop/hadoop_command_reference.htm)



26

26

## Execute a MapReduce job

- ## ❖ Run the job

**~ hadoop jar WordCount.jar /myhdfsdir/myapp/input1 /myhdfsdir/myapp/output**

- ❖ Retrieve and explore the job results

```
~ hadoop fs -copyToLocal /myhdfsdir/myapp/output1/part-r-00000 result.txt  
~ cat result.txt
```

- ❖ Clean the output HDFS directory

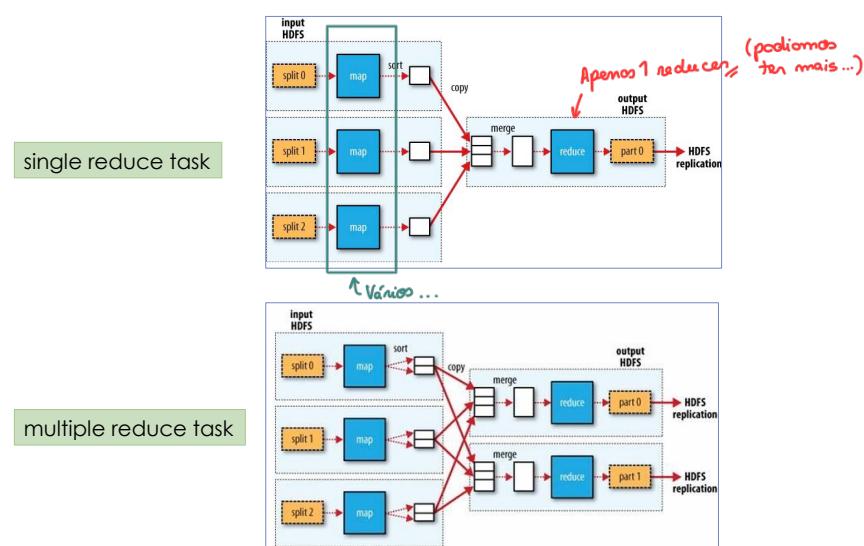
```
~ hadoop fs -rm /myhdfsdir/myapp/output1/
```



27

27

## Task execution - distinct configurations



29

29

## Java Interface

### ❖ Mapper class

- Implementation of the **map function**
- Template parameters
  - KEYIN, VALUEIN** – types of input key-value pairs
  - KEYOUT, VALUEOUT** – types of intermediate key-value pairs
- Intermediate pairs are emitted via **context.write(k, v)**

```
Api do Hadoop
class MyMapper extends Mapper<KEYIN, VALUEIN, KEYOUT, VALUEOUT> {
    @Override
    public void map(KEYIN key, VALUEIN value, Context context)
        throws IOException, InterruptedException
    {
        // Implementation
    }
}
```

Como trabalha? Qual é o contexto para escrever os values?



32

32

## Java Interface

### ❖ Reducer class

- Implementation of the **reduce function**
- Template parameters
  - KEYIN, VALUEIN** – types of intermediate key-value pairs
  - KEYOUT, VALUEOUT** – types of output key-value pairs
- Output pairs are emitted via **context.write(k, v)**

```
class MyReducer extends Reducer<KEYIN, VALUEIN, KEYOUT, VALUEOUT> {
    @Override
    public void reduce(KEYIN key, Iterable<VALUEIN> values, Context context)
        throws IOException, InterruptedException
    {
        // Implementation
    }
}
```

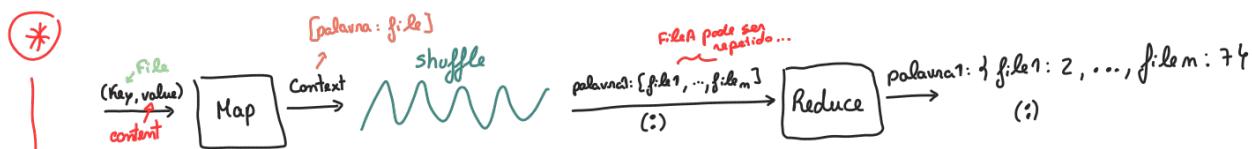
Key: [v<sub>0</sub>, v<sub>1</sub>, v<sub>2</sub>, ...]

reduce



33

33



# MapReduce – WordCount example (1)

```
public class WordCount extends Configured implements Tool {
    private static IntWritable ONE = new IntWritable(1);
    @Override
    public int run(String[] args) throws Exception {
        FileSystem fs = FileSystem.get(getConf());
        Job job = Job.getInstance(getConf()); Wrappar de String
        job.setJarByClass(WordCount.class);
        job.setJobName("WordCount");
job.setOutputKeyClass(Text.class); Wrappar a Inteiro
job.setOutputValueClass(IntWritable.class);
job.setMapperClass(Map.class); Mapper class
job.setCombinerClass(Reduce.class);
job.setReducerClass(Reduce.class); Reducer class
job.setInputFormatClass(TextInputFormat.class);
job.setOutputFormatClass(TextOutputFormat.class); input
TextInputFormat.setInputPaths(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
if (fs.exists(new Path(args[1]))) output
    fs.delete(new Path(args[1]), true);
boolean success = job.waitForCompletion(true);
return success ? 0 : 1;
    }
}
```

34

## MapReduce – WordCount example (2)

```
...  
    public static void main(String[] args) throws Exception {  
        int ret = ToolRunner.run(new WordCount(), args);  
        System.exit(ret);  
    }  
  
    public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {  
        @Override  
        public void map(LongWritable key, Text value, Context context) throws  
            IOException, InterruptedException {  
            String line = value.toString();  
            String[] words = line.split("[\t\n.,;?!-/\\[\\]\\\"']+");  
            for (String word : words) {  
                if (word.trim().length() > 0) {  
                    Text text = new Text();  
                    text.set(word);  
                    context.write(text, ONE);  
                }  
            }  
        }  
    } // end class Map
```

*Não trabalhamos diretamente com o filename do picarino*

*Wrapper de String*

*filename! = new Text(filePath...)*

*Key = new IntWritable(1);*  
*↳ definido em cima ...*

Contador!

Quase sempre  
só trabalhamos  
com o Value da  
entrada !!

⇒ Nesta caso:  
→ value é a limha  
do ficheiro!  
→ Key of

35

```

Run | Debug
public static void main(String[] args) throws Exception {
    int ret = ToolRunner.run(new InvertedIndex(), args);
    System.exit(ret);
}

public static class Map extends Mapper<LongWritable, Text, Text, Text> {
    @Override
    public void map(LongWritable key, Text value, Context context) throws
        IOException, InterruptedException {
        String line = value.toString();
        String[] words = line.split("(\\t|\\n|\\r|\\f|\\s|\\s+|\\s+\\s+)+");
        String documentId = ((org.apache.hadoop.mapreduce.lib.input.FileSplit) context.getInputSplit()).getPath().getName(); "mapreduce";
        for (String word : words) {
            if (word.trim().length() > 0) {
                Text text = new Text();
                text.set(word);
                Text docId = new Text();
                docId.set(documentId);
                context.write(text, docId);
            }
        }
    }
} // end class Map

public static class Reduce extends Reducer<Text, Text, Text, Text> {
    public void reduce(Text key, Iterable<Text> values, Context context)
        throws IOException, InterruptedException {
        java.util.Map<String, Integer> docCount = new HashMap<String, Integer>();
        for (Text val : values) {
            if (docCount.containsKey(val.toString())) {
                docCount.put(val.toString(), docCount.get(val.toString()) + 1);
            } else {
                docCount.put(val.toString(), 1);
            }
        }
        context.write(key, new Text(docCount.toString()));
    }
} // end class Reduce

```

Name do ficheiro

p1: file1  
p2: file2

p1: [file1, file1, ...]

InputFormat → Mapper → Shuffling - Sorting → Reducer → Output Format

## MapReduce – WordCount example (3)

```

...
public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        context.write(key, new IntWritable(sum));
    }
} // end class Reduce
} // end class WordCount
}
}

// Contador de palavras!

```

*Reducer*

*sum += val.get()*

*Contador de palavras!*



36

36

## MapReduce – Run the example

```

Compile
> $ javac -classpath $CLASSPATH -d WordCountDir WordCount.java

Create JAR
> $ jar -cvf WordCount.jar -C WordCountDir/ .

Run WordCount App
> $ hadoop jar WordCount.jar /in/test.txt /out/wc

See the Result
> $ hadoop fs -cat /out/wc/part-r-00000

ABOUT 1
ACCOUNT 2
ACTUAL 1
ADDITIONAL 1
ADVANCING 2
...

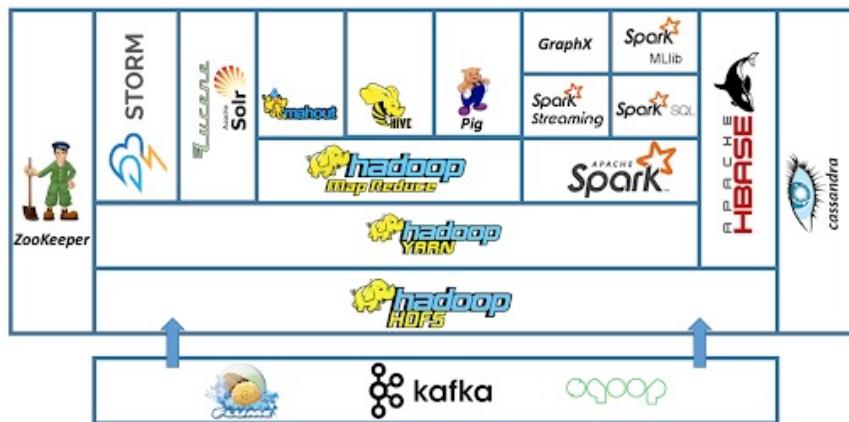
```



37

37

## Hadoop ecosystem



38

38

## Summary

- ❖ Concept
- ❖ MapReduce framework
  - Programming Model
  - Execution Plan
  - Dataflow
- ❖ MapReduce Applications
- ❖ Apache Hadoop
  - HDFS
  - MapReduce Programming

39

39